



UPPSALA
UNIVERSITET

LLM-Generated Text in Swedish: Linguistic Features and Detection Methods

Gina Welsh

Uppsala University
Department of Linguistics and Philology
Master Programme in Language Technology
Master's Thesis in Language Technology, 30 ECTS credits

February 27, 2026

Supervisor:
Birger Moëll, Uppsala University

Abstract

This thesis investigates investigate systematic differences between text generated by large language models (LLMs) and human-written text in the Swedish language. In addition, the thesis analyses significant differences in the ability to detect LLM-generated Swedish text between humans and AI-detection tools. The study will include both a formal and informal text component, that is, linguistic analyses and AI-detection experiments will be performed on two sources of data from formal and informal domains.

Contents

Preface	4
1 Introduction	5
1.1 Research questions	5
1.2 The LLMs of focus for this thesis	5
1.3 The AI-detection method of focus for this thesis	5
2 Background	6
2.1 Large Language Models and text generation	6
2.1.1 Large Language Model architectural development	6
2.1.2 LLMs for the English language	6
2.1.3 LLMs for the Swedish language	6
2.2 Textual comparison of LLM-generated and human-written text	6
2.2.1 Linguistic properties of LLM-generated text	7
2.2.2 Other properties of LLM-generated text	7
2.3 AI-generated text detection studies	8
2.3.1 AI-detection tool development	8
2.3.2 Human participant AI-detection studies	8
2.4 Background summary	8
3 Methodology	9
4 Results	10
5 Discussion	11
6 Conclusion	12

Preface

I would like to thank the following people...

1 Introduction

The intended purpose for this study is to investigate systematic differences (if any) between text generated by large language models (LLMs) and text written by humans in the Swedish language. In addition, the study will investigate whether there are any significant differences in the ability to detect LLM-generated Swedish text between humans and LLM tools. The study will include both a formal and informal text component, that is, linguistic analyses and AI-detection experiments will be performed on two sources of data from formal and informal domains.

1.1 Research questions

1. Are there systematic differences between LLM-generated and human-written formal text in Swedish?
2. Are there systematic differences between LLM-generated and human-written informal text in Swedish?
3. Are there differences between humans and AI-based detection tools in their ability to detect LLM-generated Swedish text?

1.2 The LLMs of focus for this thesis

1. OpenAI GPT-5
2. Google Gemini 3 Flash
3. GPT-SW3

1.3 The AI-detection method of focus for this thesis

- GPTZero

2 Background

This chapter will provide an overview of previous studies related to this topic. First, it will outline the development of large language models and their text generation abilities for English and Swedish. Second, it will describe previous research comparing AI-generated text with human-written text. The chapter will then outline tools and human participant studies related to AI-generated text detection.

2.1 Large Language Models and text generation

2.1.1 Large Language Model architectural development

At the time of writing this thesis, the dominant underlying architecture of LLMs is the Transformer architecture, first described in Vaswani et al. (2017). In this study, the authors designed an encoder-decoder sequence transduction model to create a network architecture that demonstrated superior performance and decreased training time in comparison to its recurrent neural network predecessors.

Specifically look at LLM architecture and text generation itself, its relationship with training on data.

2.1.2 LLMs for the English language

English has a vast amount of data - list major datasets used for LLM development in English.

Talk about compositions of English-dominant LLMs - data type

2.1.3 LLMs for the Swedish language

Ohman et al. (2023) curated a 1.2 terabyte dataset of text in all Scandinavian languages in addition to high-quality English data, called the Nordic Pile dataset. Holmstrom et al. (2023) initiated an investigation of the extent to which pre-training GPT-3 LLMs on Swedish only is advantageous over leveraging pre-trained knowledge (from English). It also examined whether English-based pre-trained models demonstrated unwanted transfer of Anglophone cultural aspects over to the Swedish context. The findings of Holmstrom et al. (2023) demonstrated that LLMs exposed to limited Swedish during training can be highly capable and transfer competencies from English "off-the-shelf", including mathematical reasoning and Swedish cultural adaptation. In addition, the study showed that GPT-3 models did not appear to exhibit unwanted over-generalisation of cultural aspects to Swedish.

2.2 Textual comparison of LLM-generated and human-written text

Several studies have already done that analyse the linguistic characteristics of LLM-generated text in English.

2.2.1 Linguistic properties of LLM-generated text

Several studies have already done that analyse the linguistic characteristics of LLM-generated text in English.

Munoz-Ortiz et al. (2024) compared the linguistic components of human and LLM-generated news articles and compared their patterns among six generative LLMs. For their prompting strategy, they asked the models to generate a news paragraph based on the news article headline its first words. For their linguistic analysis, they compared the human-written and LLM-generated text groups through differences in their linguistic structures, including vocabulary distribution, part-of-speech tag distributions, dependency arc lengths, constituent lengths, cosine similarity, and emotional category tagging. The results of this analysis showed significant differences between the human-written and LLM-generated text on all textual analysis metrics. At the lexical level, 5 of the 6 LLM models outputted text with a more restricted vocabulary than the human-written text. For semantics, the models were found to have less of an inclination to display aggressive negative emotions (such as fear or anger) than humans.

Liao et al. (2023) analysed the differences between human-written and ChatGPT-generated text within the medical domain. Medical texts were compared by their vocabulary distribution, part-of-speech tag distribution, dependency arcs, sentiment analysis, and perplexity. In addition, they implemented machine learning methods to detect ChatGPT-generated medical text. The results of the study found that human-written medical texts contained more concrete information, had higher vocabulary diversity, and contained information that human experts would consider more useful. The ChatGPT-generated medical texts appeared more fluent and logical, but expressed more generalised terminology rather than specific, effective information for the problem's context. Overall, this would point to the human-written medical text being more tailored to the situational modelling of the medical case (situational modelling defined in Mahowald et al., 2024) rather than the LLM-generated text.

- discuss differences in text stylistic

2.2.2 Other properties of LLM-generated text

Previous work has also shown evidence of a distinction between the reasoning patterns of LLM tools when compared to those of humans. (CITE) evaluated the medical reasoning capabilities of DeepSeek R1, computing its outputs to the reasoning patterns of medical domain human experts. The study performed qualitative and quantitative analyses of AI- and human-medical reasoning for 100 clinical cases from the MedQA dataset. The study found a 93% diagnostic accuracy for DeepSeek R1's performance. Many of the LLM's errors were attributed to over-focus on the details of a medical problem, rather than a more holistic and intuitive (?) understanding by humans, as well as an insufficient understanding of medical protocols when compared to human experts. The study also found a strong correlation between the length of output and the error rate of LLMs; that is, results from LLMs that were longer were more likely to contain incorrect diagnoses or medical details.

- discuss differences in text stylistic

2.3 AI-generated text detection studies

2.3.1 AI-detection tool development

2.3.2 Human participant AI-detection studies

Landberg (2024) investigated the effectiveness of AI chatbots to detect texts created by AI in Swedish and English. The main component of the study A recent study investigated LLM reasoning in Swedish. Moell and Boye (2025) studied the performance of LLMs on language complexity measurement tasks (LIX readability metric and Average Dependency Distance). The results of the study found that all models demonstrated at least some capacity for each of these tasks, but that LLMs remain prone to errors in explicit calculations and syntactic analysis.

2.4 Background summary

3 Methodology

text here

4 Results

text here

5 Discussion

text here

6 Conclusion

text here

Bibliography

Vaswani, Ashish, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin (2017). “Attention Is All You Need”. *CoRR* abs/1706.03762. arXiv: [1706.03762](https://arxiv.org/abs/1706.03762). URL: <http://arxiv.org/abs/1706.03762>.